

# Clojure & DynamoDB

2013-07-08

Stephen Johnston

@stevejjohnston



# DynamoDB Overview

- ✦ In 2007, Amazon published “Dynamo: Amazon’s Highly Available Key-value store”
- ✦ The technology behind Amazon S3.
- ✦ Several open-source implementations released based on this paper: Riak, Project Voldemort
- ✦ In 2012, Amazon offers “DynamoDB”, a nosql database as a service
- ✦ Tunable read/write performance. (Provisioned capacity.)



# Data Model

- ✦ DynamoDB allows users to setup one or more tables
- ✦ Each table has a primary key: Either a 'hash' or a 'hash + range'
- ✦ The remaining columns are schema-less
  - ✦ Each row in the table can have a varying number of columns
- ✦ Columns can be either string, numeric, or binary



# Quick note on hashing

- ✦ The hash key determines where the data is written.
- ✦ We know that Dynamodb scales out horizontally.
- ✦ Thus, it is important that hashes have a nice even distribution.



# Features

- ✦ Replicates data across 3 availability zones (data centers within the same region)
- ✦ Single digit millisecond response for get/put requests
- ✦ No administration
- ✦ Scales throughput / storage
- ✦ Local secondary indexes
- ✦ Transactions



# Pricing

- ✦ Free Tier
  - ✦ 5 writes per second, 10 reads per second, 100MB data.
- ✦ After free tier, storage is 25¢ per GB per month
- ✦ Writes are \$0.0065 per hour for each 10 units of write capacity
- ✦ Reads are \$0.0065 per hour for each 50 units of read capacity
- ✦ Pay as you go



# Pricing example

- 500 reads per second, 500 writes per second, 50GB of data
  - $500 \text{ read} * \$0.0065 / 50 \text{ read} * 24 * 365 / 12 = \$47$
  - $500 \text{ writes} * \$0.0065 / 10 \text{ writes} * 24 * 365 / 12 = \$237$
  - $50 * 25\text{¢} = \$12.50$
  - Grand total per month: \$296.50



# Limits

- Reads and writes are capped at your provisioned capacity
- Your provisioned capacity can be increased as often as you like, but can only be decreased 4 times a day
- Can only increase capacity 100% at a time
- Tables have no practical limit on number of rows or bytes
- 256 tables per region
- An item (or row) has a max size of 64kb
- Batch reads: 100 rows at a time, or 1MB of data.
- Batch writes: 25 rows at a time, or 1MB of data.



# DynamoDB and Clojure

- ✦ Weavejester's library: rotary
- ✦ ptaoussanis' library: faraday
- ✦ Clojure maps are a nice fit for schemaless items in DynamoDB
- ✦ { :user-id 100 :first-name "Bob" :last-name "Smith" }



# Sample use case:

## 3 Billion phone calls a day

- ✦ Needs 35,000 writes per second
- ✦ Need to be able to reverse lookup a phone number
- ✦ Need to be able to see all the calls for a phone number, sorted by calltime (or sorted by another column, like duration or number-dialed.)
- ✦ Estimated Price?
  - ✦ Reserved capacity price: ~ \$8,000 per month



# Table setup

- my-table (reverse lookup table)
  - Primary hash key: 'phonenumber' as string
  - Other columns: 'name' as string
  - Could grow the table to include additional columns as needed (address, account #, etc)
- phone-metadata
  - Hash key: 'phonenumber' as string
  - Range key: 'calltime' as number
  - Other columns: duration, latitude, longitude, numberdialed



Demo ptaoussanis/faraday